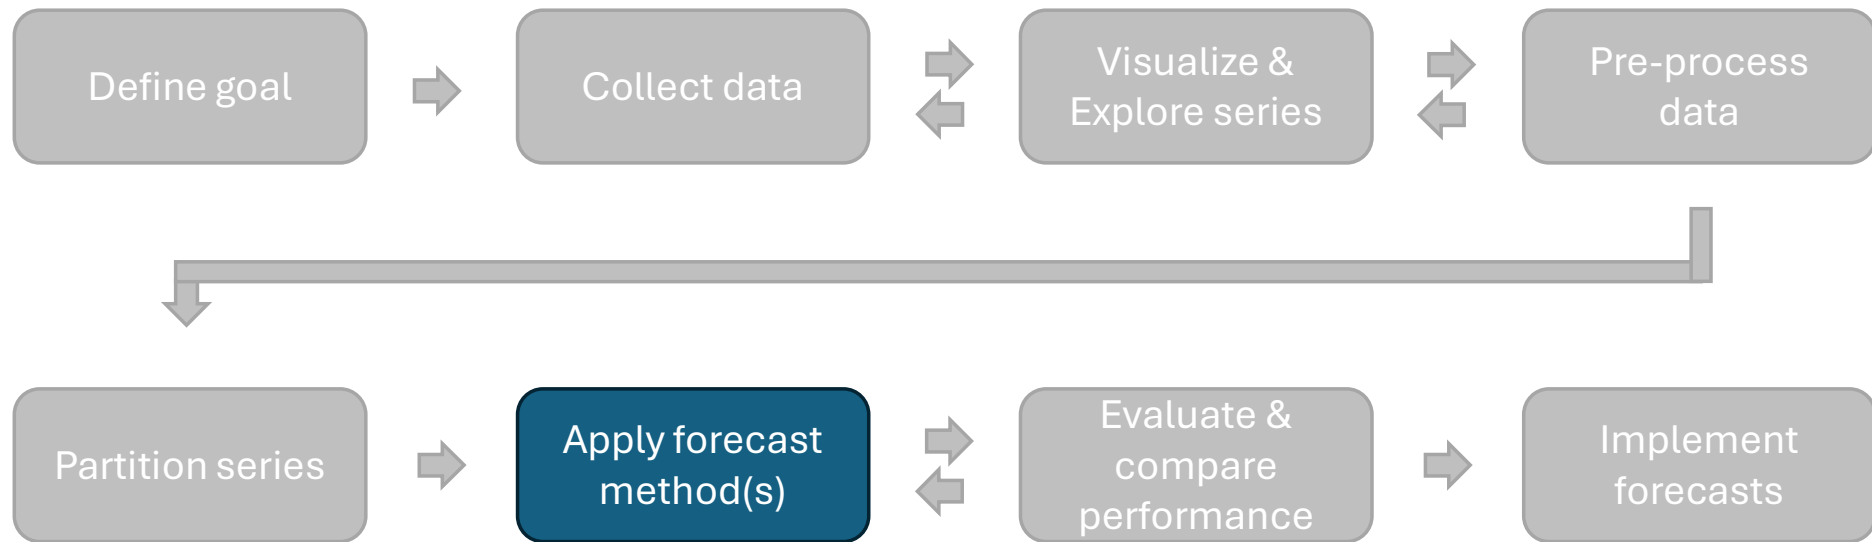


# Time series Forecasting

## Part 2: Smoothing Methods

# The forecasting process



# Components of a Time Series - Recap

1. **Level** (always present)
2. **Trend**: steady increase/decrease over time.
3. **Seasonality**: pattern that repeats itself every season
4. Random **noise** (always present)

# Reminder: Modeling Principles

## Time series Analysis

- Reasonableness and parsimony

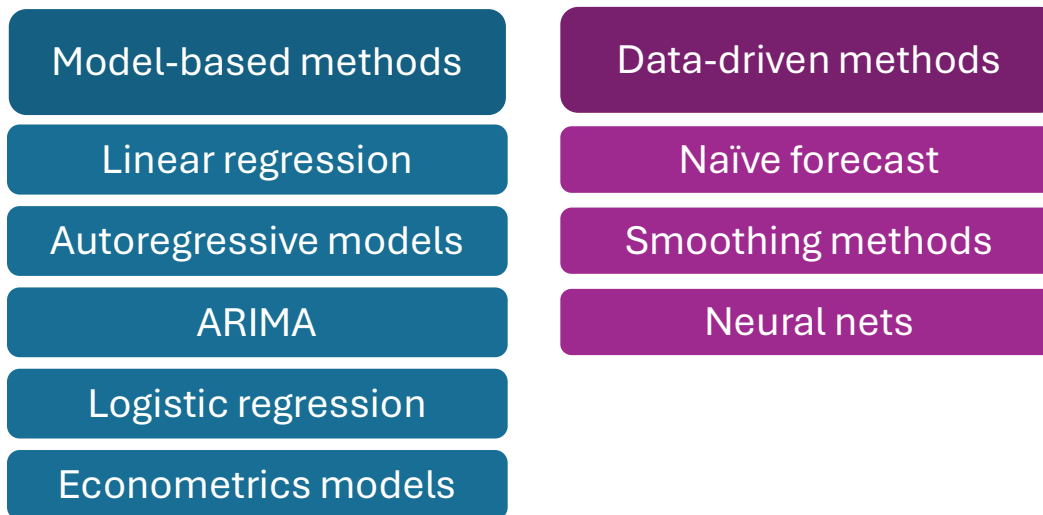
- Goodness of fit (residual analysis)

## Time series Forecasting

- Forecast accuracy

- Parsimony and reasonableness

# Why so many different methods?



# Smoothing methods

**Moving average**

**Exponential  
smoothing**

- Smoothing methods are useful for
  - Data visualization
  - Removing seasonality and computing seasonal indexes
  - Forecasting

# The Moving Average Method

## **The Idea:**

Forecast future points by using an average of several past points

## **Uses:**

Time series visualisation  
Computing seasonal indexes  
Forecasting

## **Advantages:**

Simple, popular

## **Disadvantages:**

Forecast only in series that lack seasonality and trend

## **Key concept:**

Width of window

# Notation

$t = 1, 2, 3, \dots$

An index denoting the time period of interest.  
 $t = 1$  is the first period in a series.

$y_1, y_2, y_3, \dots, y_n$

A series of  $n$  values measured over  $n$  time periods,  
where  $y_t$  denotes the series value at time  $t$ .

$F_t$

The forecasted value for time period  $t$ .

$F_{t+k}$

The  $k$ -step-ahead forecast.

$e_t$


The forecast error for time period  $t$ , which is the difference between the actual value and the forecast at time  $t$ , and equal to  $y_t - F_t$ .



# Two types of windows


## Centered moving average:

based on a window centered around time  $t$



## Trailing moving average:

based on a window from time  $t$  and backwards



# Moving averages for visualizing time series

A time-plot of the **moving averages** can help reveal the LEVEL and TREND of a series, by filtering out the seasonal and random components


# Example: Ridership on Amtrak Trains



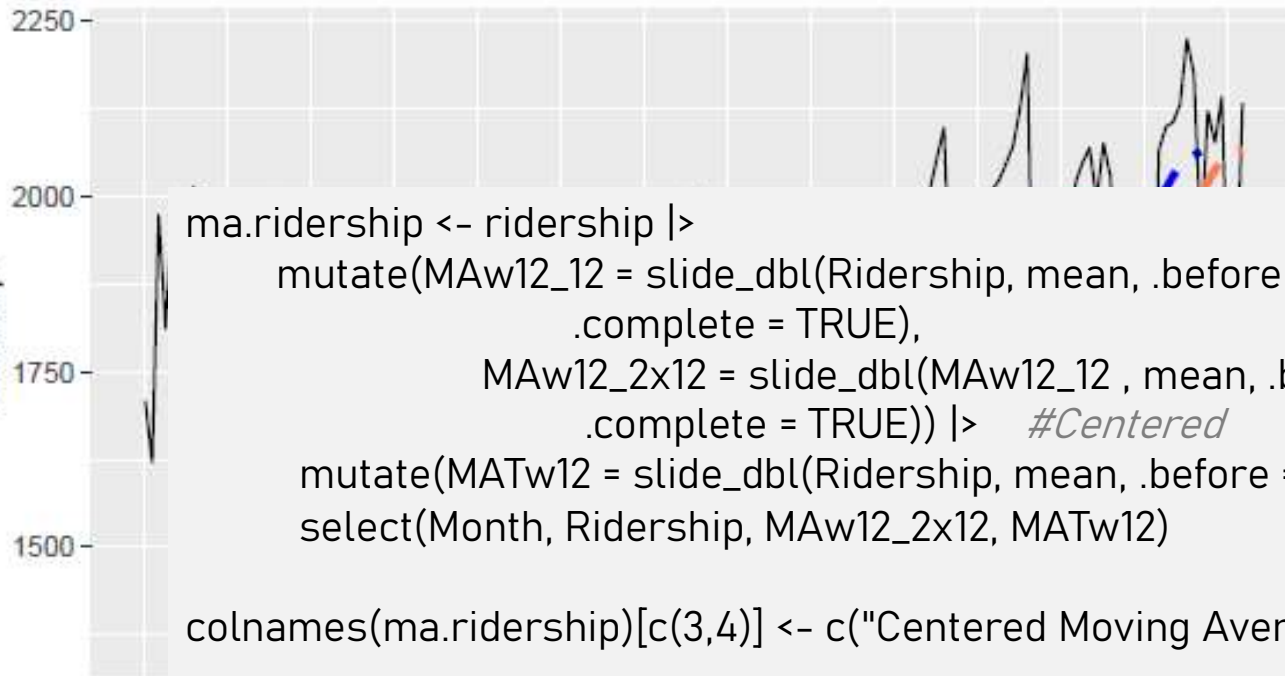
US Railway company  
Monthly ridership, Jan 1991  
to Mar 2004 (Amtrak.csv)

## Moving average

Centered moving average and trailing moving average with window width  $w=12$ , overlaid on Amtrak ridership series



## Moving average



Centered moving average and trailing moving average with window width  $w=12$ , overlaid on Amtrak ridership series

```
ma.ridership <- ridership |>
  mutate(MAw12_12 = slide_dbl(Ridership, mean, .before = 5, .after = 6,
    .complete = TRUE),
    MAw12_2x12 = slide_dbl(MAw12_12, mean, .before = 1, .after = 0,
    .complete = TRUE)) |> #Centered
  mutate(MATw12 = slide_dbl(Ridership, mean, .before = 11, .after = 0, .complete = TRUE)) |> #Trailing
  select(Month, Ridership, MAw12_2x12, MATw12)

colnames(ma.ridership)[c(3,4)] <- c("Centered Moving Average", "Trailing Moving Average")

ma.ridership <- ma.ridership |> gather(model, value, Ridership:`Trailing Moving Average`)

m ma.ridership |> ggplot(aes(x = Month, y = value)) +
  geom_line(aes(color = model, linetype = model, size = model)) +
  scale_size_manual(values = c(1.2, 0.4, 1.2)) +
  scale_linetype_manual(values = c("dotted", "solid", "longdash")) +
  theme(legend.position = "bottom") +
  scale_x_yearmonth(date_breaks = "2 years", date_labels = "%Y") +
  labs(x = "Time", y = "Ridership")
```




## Centered moving average with an even window

```
ma.ridership <- ridership |>
  mutate(MAw12_12 = slide_dbl(Ridership, mean, .before = 5, .after = 6,
    .complete = TRUE),
    MAw12_2x12 = slide_dbl(MAw12_12, mean, .before = 1, .after = 0,
    .complete = TRUE))    #Centered
```

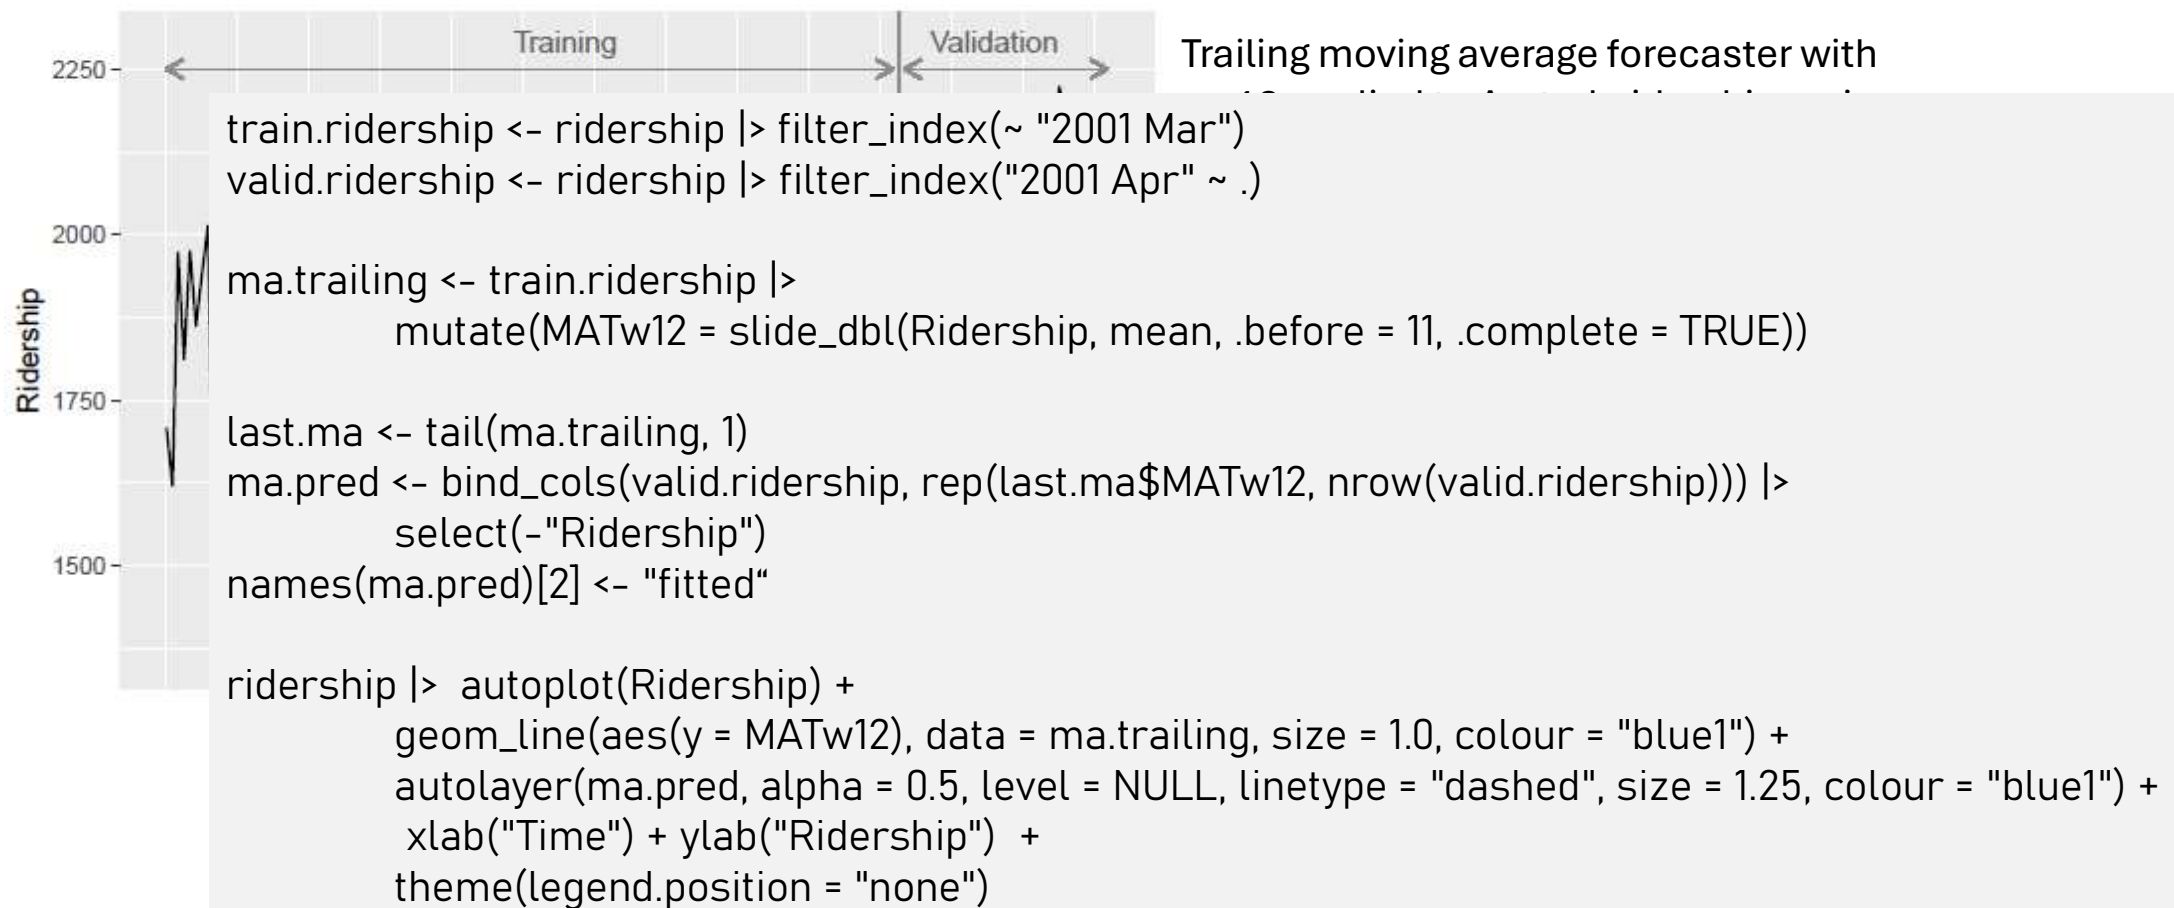
$$\begin{aligned}MAw12\_2x12 &= \frac{1}{2} \left[ \frac{1}{12} (y_{t-5} + y_{t-4} + y_{t-3} + y_{t-2} + y_{t-1} + y_t + y_{t+1} + y_{t+2} + y_{t+3} + y_{t+4} + y_{t+5} + y_{t+6}) + \right. \\ &\quad \left. \frac{1}{12} (y_{t-6} + y_{t-5} + y_{t-4} + y_{t-3} + y_{t-2} + y_{t-1} + y_t + y_{t+1} + y_{t+2} + y_{t+3} + y_{t+4} + y_{t+5}) \right] \\ &= \frac{1}{24} y_{t-6} + \frac{1}{12} y_{t-5} + \frac{1}{12} y_{t-4} + \frac{1}{12} y_{t-3} + \frac{1}{12} y_{t-2} + \frac{1}{12} y_{t-1} + \frac{1}{12} y_t + \\ &\quad \frac{1}{12} y_{t+1} + \frac{1}{12} y_{t+2} + \frac{1}{12} y_{t+3} + \frac{1}{12} y_{t+4} + \frac{1}{12} y_{t+5} + \frac{1}{24} y_{t+6}\end{aligned}$$

# The Moving Average Method



Trailing moving average forecaster with  $w=12$  applied to Amtrak ridership series.

# The Moving Average Method





# Choosing window width ( $w$ )

- Balance over- and under-smoothing
- Wide window – global trend, narrow window – reveal local trend
- If no seasonality, use a narrow window (under-smoothing)

## **Test yourself**

For a seasonal series, what window width should you use?

1. Smaller than the # seasons
2. Larger than the # seasons
3. Equal to the # seasons



## Hands-On # 2.1

- Fortified wine has the largest market share of the six types of wine. You are asked to focus on fortified wine sales alone
- Fit moving average model with a window of size 5 and size 12
- What patterns do these models suggest?
- (Hint: don't forget to split the data into training and validation sets)

# Removing trend and/or seasonality

## **Differencing:**

Taking the difference between two consecutive observations.

Usage

Removing a trend and/or seasonality from a time series

*Lag* – 1 difference:  $y_t - y_{t-1}$

Removing trend



E.g., monthly  
seasonality  $m=12$ ,  
daily seasonality  
 $m=7$

*Lag* –  $m$  difference:  $y_t - y_{t-m}$

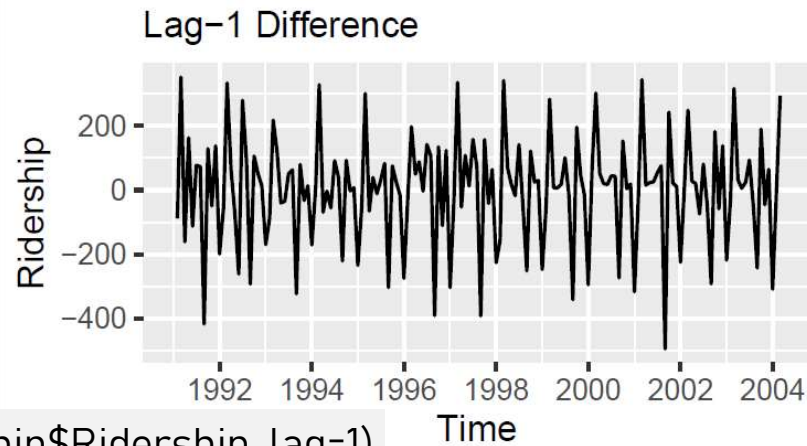
Removing seasonality with  $m$  seasons

Double-differencing: difference the differenced series


# Example: Ridership on Amtrak Trains



`diff(ridership$Ridership, lag=12)`



`diff(ridership$Ridership, lag=1)`



`diff(diff(ridership$Ridership, lag=12), lag=1)`

# Methods require no trend and/or seasonality

- Moving average
- (Simple) Exponential smoothing

# The simple exponential smoothing Method

## **The Idea:**

Forecast future points by using an exponential weighted average of several past points

## **Uses:**

Forecasting  
Automated forecasting  
Capture well for local patterns  
Visualization  
Creating seasonal indexes

## **Advantages:**

Simple to understand  
Gives more influence on recent information  
Storage/computation efficient (we only need to store the last forecast and most recent observation)

## **Disadvantages:**

Forecast only in series that lack seasonality and trend\*  
Forecast into the future & one-step-ahead forecast the same

## **Key concept:**

Smoothing constant  $\alpha$

# Exponential smoothing

## Types of exponential smoothing

Simple exponential smoothing (no trend or seasonality)

Advance exponential smoothing (trend and/or seasonality)

Holt's method (with trend but no seasonality)

Winter's method (with trend and seasonality)

# Simple exponential smoothing

- Assume that the series has only level ( $L_t$ ) and noise (unpredictable)

$$F_{t+1} = L_t$$

Forecasts are estimated as level at the most recent time

$$L_t = \alpha y_t + (1 - \alpha)L_{t-1}$$

Exp smoothing is an adaptive algorithm.

It adjust the most resent forecast (level) based on actual data

$$0 < \alpha \leq 1$$

$\alpha$  = the smoothing constant

Initialization:  $F_t = L_1 = y_1$



# Why 'exponential smoothing'?

- Let's examine  $F_{t+1} = L_t$  where  $L_t = \alpha y_t + (1 - \alpha)L_{t-1}$

$$\begin{aligned}
 F_{t+1} &= L_t \\
 &= \alpha y_t + (1 - \alpha)[\alpha y_{t-1} + (1 - \alpha)L_{t-2}] = \\
 &= \alpha y_t + \alpha(1 - \alpha)y_{t-1} + (1 - \alpha)^2 L_{t-2} = \\
 &= \alpha y_t + \alpha(1 - \alpha)y_{t-1} + \alpha(1 - \alpha)^2 y_{t-2} + \dots
 \end{aligned}$$

**Weights decrease exponentially into the past!**

Required data  
storage and  
computational  
time!

$$= F_t + \alpha e_t$$

Update **previous  
forecast**

$\alpha$  controls the  
**degree of 'leaning'**

**Active learner!**

By an amount that depends on the  
**error** in the previous forecast

See page 94

1 para

# The smoothing constant $\alpha$

- Controls the degree of ‘leaning’

$$0 \leftarrow \alpha \rightarrow 1$$

Slow learner

Fast learner

Past obs.

Have a large influence on the forecast

Have little to no influence forecasts

Over-smoothing


Under-smoothing

- Selecting  $\alpha$

Common values: 0.1 or 0.2

Trail & error: effect on visualisation

Minimise RMSE or MAPE of training data



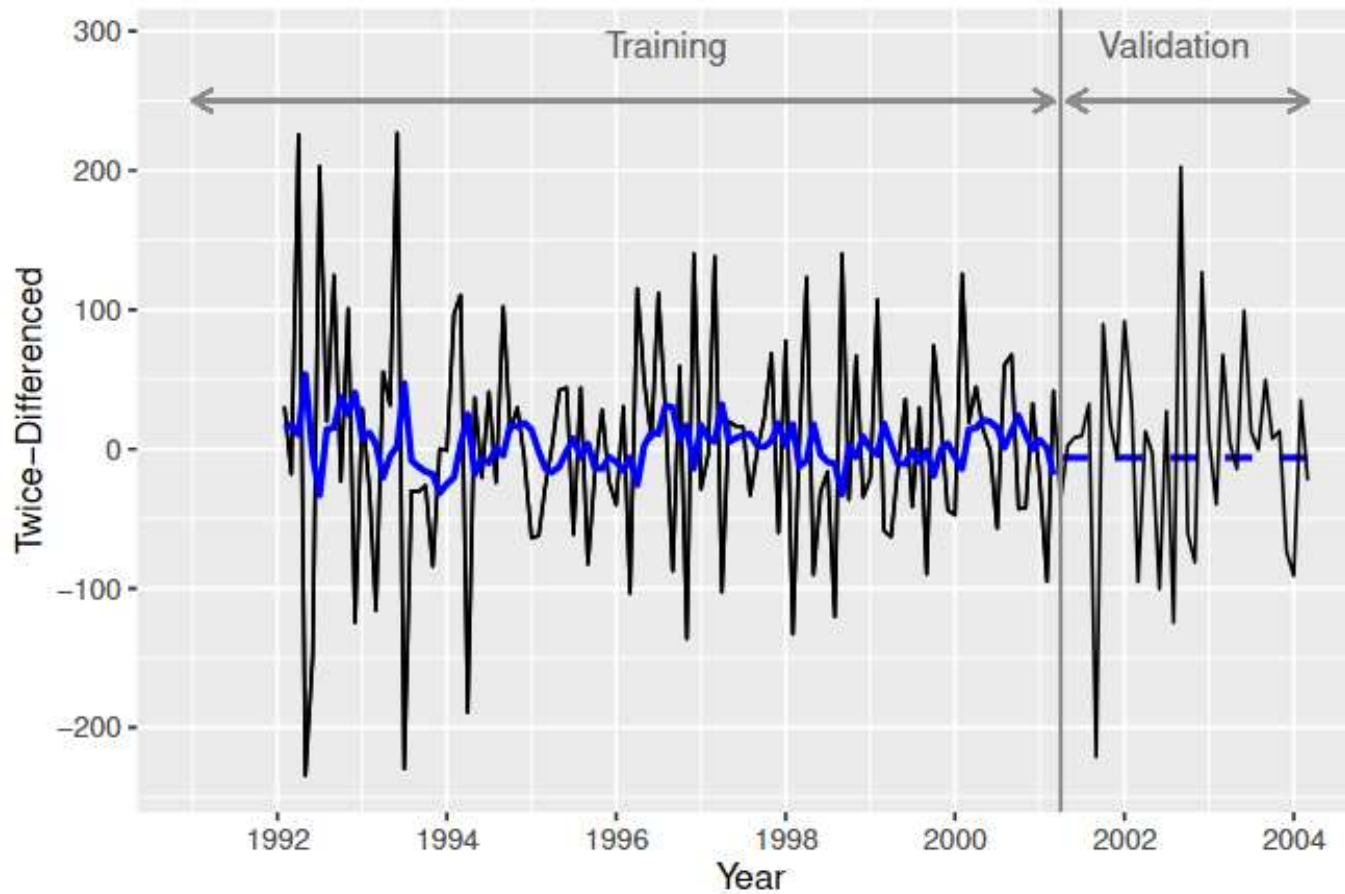
$\alpha$	$\alpha(1-\alpha)$	$\alpha(1-\alpha)^2$	$\alpha(1-\alpha)^3$
0.9	0.09	0.009	0.0009
0.5	0.25	0.125	0.0625
0.2	0.16	0.128	0.1024
0.1	0.09	0.081	0.0729

# Example: Ridership on Amtrak Trains



US Railway company  
Monthly ridership, Jan 1991  
to Mar 2004 (Amtrak.csv)

# Exponential smoothing

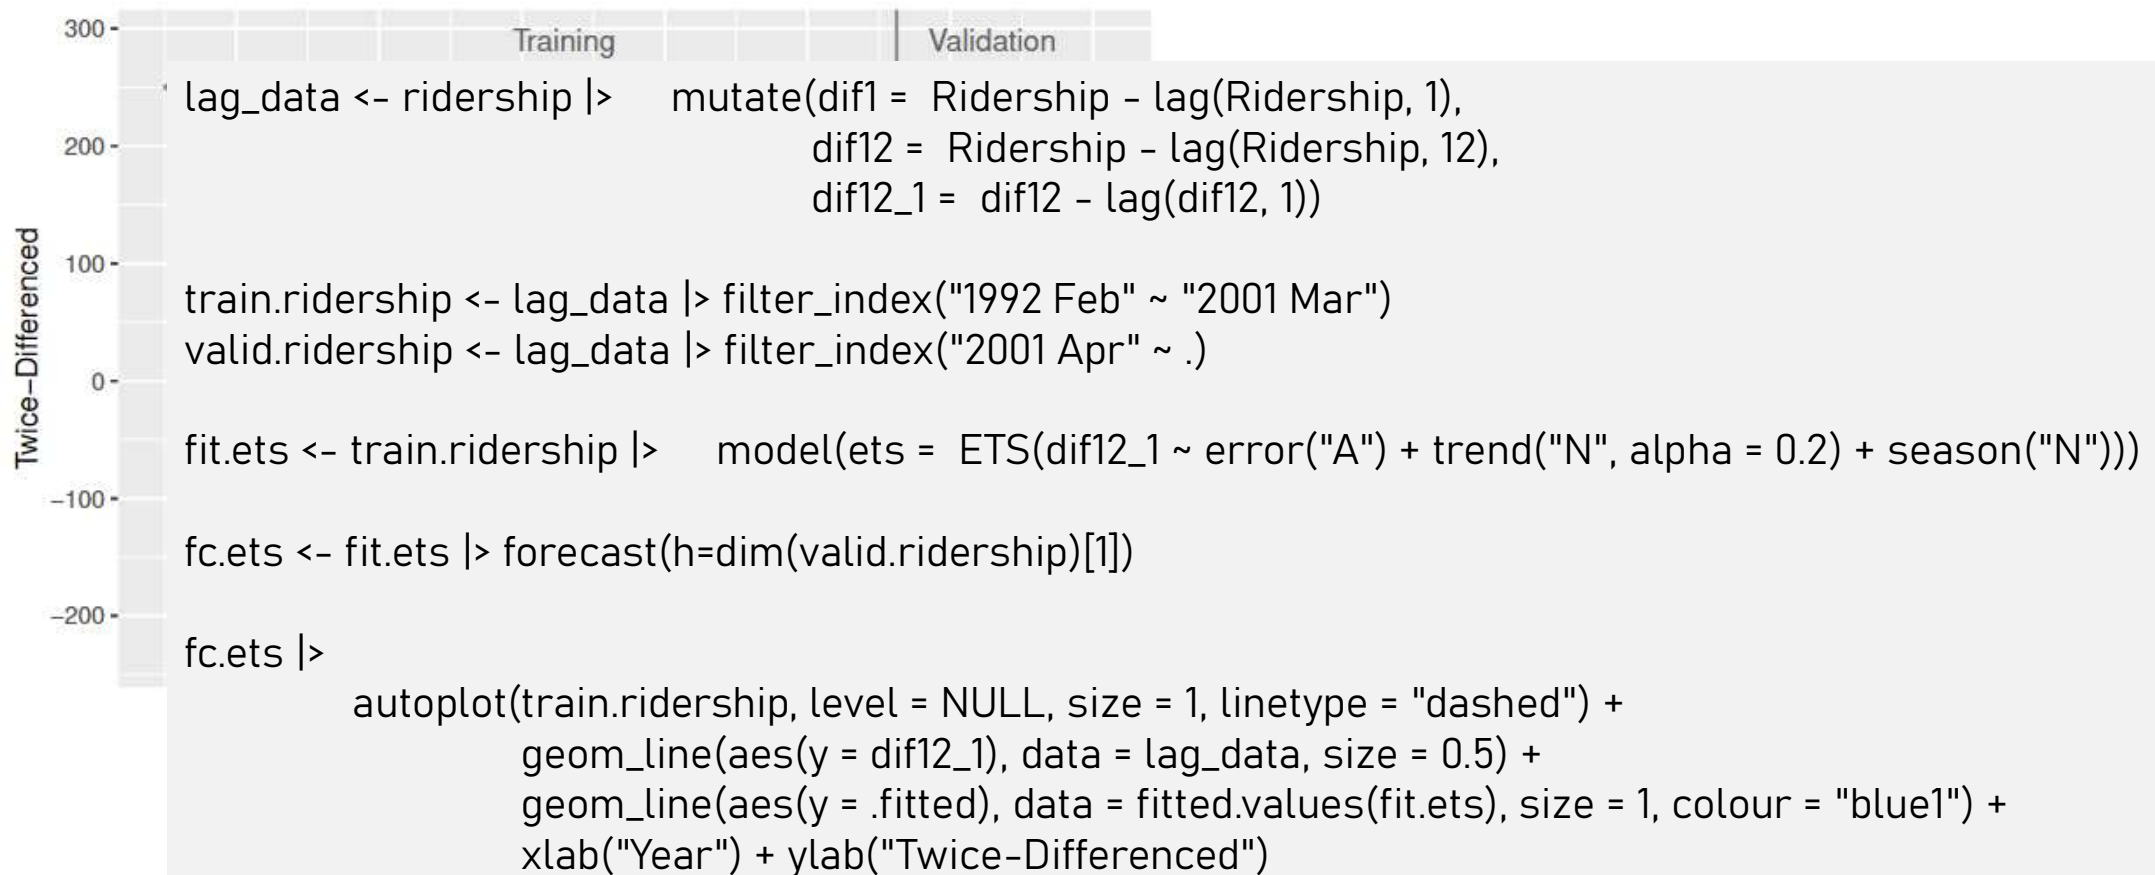


Exp Smoothing ( $\alpha=0.2$ ) applied to twice-differenced Amtrak series



# Exponential smoothing

Hands on activity –  
generate auto ETS



# Some useful commands

- *report*(fit.ets)

```
> report(fit.ets)
Series: dif12_1
Model: ETS(A,N,N)
  Smoothing parameters:
    alpha = 0.2

  Initial states:
    1[0]
    14.35123

  sigma^2: 8401.487

      AIC      AICc      BIC
1513.012 1513.125 1518.413
```

# Some useful commands

- Forecasted values for the training data
- `pred.values.ets <- fitted.values(fit.ets)`
- `View(pred.values.ets)`

	.model	Month	.fitted
1	ets	1992 Feb	14.35123000
2	ets	1992 Mar	17.59938400
3	ets	1992 Apr	10.48070720
4	ets	1992 May	53.54616576
5	ets	1992 Jun	-4.07626739
6	ets	1992 Jul	-33.07381391
7	ets	1992 Aug	14.09354887
8	ets	1992 Sep	15.27463910
9	ets	1992 Oct	37.17811128
10	ets	1992 Nov	25.10968902
11	ets	1992 Dec	40.25935122
12	ets	1993 Jan	7.27248097
13	ets	1993 Feb	11.65258478
14	ets	1993 Mar	3.52506782
15	ets	1993 Apr	-20.37214574
16	ets	1993 May	-5.18251659
17	ets	1993 Jun	2.13578673
18	ets	1993 Jul	47.08562938
19	ets	1993 Aug	-8.26269650

# Some useful commands

- Forecasted values for the validation data
- `fc.ets`

```
> fc.ets
# A fable: 36 x 4 [1M]
# Key:      .model [1]
#   .model      Month      dif12_1 .mean
#   <chr>      <mth>      <dist> <dbl>
1 ets      2001 Apr  N(-6.3, 8401) -6.35
2 ets      2001 May  N(-6.3, 8738) -6.35
3 ets      2001 Jun  N(-6.3, 9074) -6.35
4 ets      2001 Jul   N(-6.3, 9410) -6.35
5 ets      2001 Aug   N(-6.3, 9746) -6.35
6 ets      2001 Sep   N(-6.3, 10082) -6.35
7 ets      2001 Oct   N(-6.3, 10418) -6.35
8 ets      2001 Nov   N(-6.3, 10754) -6.35
9 ets      2001 Dec   N(-6.3, 11090) -6.35
10 ets     2002 Jan   N(-6.3, 11426) -6.35
# i 26 more rows
# i Use `print(n = ...)` to see more rows
> |
```



# Exponential smoothing – link to moving average

Moving average with  $w = \frac{2}{\alpha} - 1$  will provide (approximately) equal result to the result of exponential smoothing with smoothing constant  $\alpha$ .

# Advanced exponential smoothing

Double exponential smoothing, Holt's method. **Additive trend**

$$F_{t+k} = L_t + kT_t$$

Assume a trend that can change over time (local trend). The level changes from one period to next by a fix amount

$$L_t = \alpha y_t + (1 - \alpha)(L_{t-1} + T_{t-1})$$

Level estimate at time  $t$ , a weighted average of observation at time  $t$  and the level in the previous period adjusted for trend

$$T_t = \beta(L_t - L_{t-1}) + (1 - \beta)T_{t-1}$$

Trend estimate at time  $t$ , a weighted average of the most recent info on the change in level and the trend in the previous period

$$0 < \alpha \leq 1$$

$$0 < \beta \leq 1$$

Smoothing constants determine the rate of learning. As closer both to 1 as faster the learning (more weight to more recent obs.)

# Advanced exponential smoothing

Two types of errors (additive trend)

*Additive error*

$$y_t = L_t + T_t + e_t$$

The error has a similar magnitude irrespective of the current level and/or trend of the series

Multiplicative error

$$y_t = (L_t + T_t)(1 + e_t)$$

The error proportionally increases (decreases) with level and/or trend changes.

# Advanced exponential smoothing

## Multiplicative trend

$$F_{t+1} = L_t \times T_t^k$$

Assume a trend that can change over time (local trend). The level changes from one period to the next by a ~~fix amount~~ factor

$$L_t = \alpha y_t + (1 - \alpha)(L_{t-1} \times T_{t-1})$$

Updating equation for level at time  $t$

$$T_t = \beta(L_t/L_{t-1}) + (1 - \beta)T_{t-1}$$

Updating equation for trend at time  $t$

$$0 < \alpha \leq 1$$

$$0 < \beta \leq 1$$

Smoothing constants determine the rate of learning. As closer both to 1 as faster the learning (more weight to more recent obs.)

# Advanced exponential smoothing

Two types of errors (Multiplicative trend)

## *Additive error*

$$y_t = L_t \times T_t + e_t$$

The error has a similar magnitude irrespective of the current level and/or trend of the series

## Multiplicative error

$$y_t = (L_t \times T_t)(1 + e_t)$$

The error proportionally increases (decrease) with level and/or trend changes.

# Advanced exponential smoothing

Holt-Winter's exponential smoothing. **Additive seasonality**

$$F_{t+k} = L_t + kT_t + S_{t+k-m}$$

Assume a trend that can change over time (local trend). The level changes from one period to the next by a fixed amount. The values of different seasons differ by a fixed amount

$$L_t = \alpha(\underbrace{y_t - S_{t-m}}_{\text{Seasonal adjusted value}}) + (1 - \alpha)(L_{t-1} + T_{t-1})$$

Seasonal adjusted value

Updating equation for level at time  $t$

$$T_t = \beta(L_t - L_{t-1}) + (1 - \beta)T_{t-1}$$

Updating equation for trend at time  $t$

$$S_t = \gamma(\underbrace{y_t - L_t}_{\text{Level adjusted value}}) + (1 - \gamma)S_{t-m}$$

Level adjusted value

Updating equation for seasonality at time  $t$

$$0 < \alpha, \beta, \gamma \leq 1$$

Smoothing constants determine the rate of learning.

# Advanced exponential smoothing

Holt-Winter's exponential smoothing. **Multiplicative seasonality**

$$F_{t+k} = (L_t + kT_t) \times S_{t+k-m}$$

Assume a trend that can change over time (local trend). The level changes from one period to the next by a fixed amount. The values of different seasons differ by percentage amount

$$L_t = \alpha(\underbrace{y_t / S_{t-m}}_{\text{Seasonal adjusted value}}) + (1 - \alpha)(L_{t-1} + T_{t-1})$$

Seasonal adjusted value

Updating equation for level at time  $t$

$$T_t = \beta(L_t - L_{t-1}) + (1 - \beta)T_{t-1}$$

Updating equation for trend at time  $t$

$$S_t = \gamma(\underbrace{y_t / L_t}_{\text{Level adjusted value}}) + (1 - \gamma)S_{t-m}$$

Level adjusted value

Updating equation for seasonality at time  $t$

$$0 < \alpha, \beta, \gamma \leq 1$$

Smoothing constants determine the rate of learning.




# Example: Ridership on Amtrak Trains



US Railway company  
Monthly ridership, Jan 1991  
to Mar 2004 (Amtrak.csv)




# Advanced exponential smoothing



Forecasts from the Holt-Winter's exponential smoothing applied to the Amtrak ridership series.

# Advanced exponential smoothing



Forecasts from the Holt-Winter's exponential smoothing applied to the Amtrak ridership series.

```
fit.ets.HW <- train.ridership |>
  model(ets = ETS(Ridership ~ error("M") + trend("A") + season("A")))

fc.ets.HW <- fit.ets.HW |> forecast(h = nrow(valid.ridership))

ridership |> autoplot(Ridership) +
  geom_line(aes(y = .mean), data = fc.ets.HW, colour = "blue1", linetype = "dashed") +
  autolayer(fitted.values(fit.ets.HW), alpha = 0.5, level = NULL, colour = "blue1") +
  xlab("Time") + ylab("Ridership")
```

# Advanced exponential smoothing

```
> report(fit.ets.HW)
Series: Ridership
Model: ETS(M,A,A) error("M") + trend("A") + season("A"))

Smoothing parameters:
  alpha = 0.5517889 Estimate far from zero => level adopts locally
  beta  = 0.0001119929 Estimate close zero => trend & seasonality global
  gamma = 0.0001186167

Initial states:
  l[0]      b[0]      s[0]      s[-1]      s[-2]      s[-3]      s[-4]      s[-5]      s[-6]
  1838.577  0.8081064  28.44536  -11.21868  -0.533687  -124.2766  200.1969  146.8103  36.59075
  s[-7]      s[-8]      s[-9]      s[-10]     s[-11]
  76.04396   60.15303   44.17906  -249.4682  -206.9221

sigma^2: 0.0012

AIC      AICc      BIC
1617.596  1623.424  1665.403
```

# Advanced exponential smoothing

# Final states:

```
> n <- nrow(components(fit.ets.HW))
```

```
> components(fit.ets.HW)[n, c("level", "slope")] # level and trend
```

# A tibble: 1 × 2

level	slope
<dbl>	<dbl>1
1945.	0.810

```
> t(components(fit.ets.HW)[(n-11):n, c("season")]) # s1 to s12
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]
season	60.13771	76.05659	36.58921	146.8106	200.1936	-124.2748	-0.5324645	-11.21618	28.44364
	[,10]	[,11]	[,12]						
season	-206.9249	-249.471	44.1899						

# Advanced exponential smoothing

The ETS() function

- Provide 18 model choices, combination of:

Trend	Seasonality		
	None	Additive	Multiplicative
None	ZNN	ZNA	ZNM
Additive	ZAN	ZAA	ZAM
Multiplicative	ZMN	ZMA	ZMM

where Z can be set to A or M corresponding to additive or multiplicative error, respectively.

```
train.ridership |> model(ets = ETS(Ridership ~ error("M") + trend("A") + season("A")))
```

# Automated exponential smoothing

- If you do not know which model to choose, you can use the ETS function to do automated model selection.
- It will fit all the models and select the one that minimised *Akaike's Information Criterion* (AIC),
- AIC combines fit to the training data with a penalty for the number of smoothing parameters.
- The AMM, AAM, AMA, and MMA models will not be considered as they can lead to numerical difficulties.

Trend	Seasonality		
	None	Additive	Multiplicative
None	ZNN	ZNA	ZNM
Additive	ZAN	ZAA	A → ZAM
Multiplicative	ZMN	ZMA	M → ZMM A → ZMM

# Automated exponential smoothing

- To include the restricted models when searching for a model, include the options *restrict = FALSE* in the ETS function.
- You can automatically select a model from a subset of pre-specified models.
- For example,  
    `train.ridership |> model(ets= ETS(Ridership ~ error("M"))`  
will consider all the models with multiplicative error except the MMA model (unless you include *restrict = FALSE* ).

# Automated exponential smoothing - example

```
fit.ets.auto <- train.ridership |> model(ets = ETS(Ridership))
```

```
pred.values.ets.auto <- fitted.values(fit.ets.auto)
```

```
fc.ets.auto <- fit.ets.auto |> forecast(h=dim(valid.ridership)[1])
```

*error("M"),*  
*trend("N"),*  
*season("A")*

```
> report(fit.ets.auto)
Series: Ridership
Model: ETS(M,N,A)
Smoothing parameters:
  alpha = 0.5550234
  gamma = 0.0001000004

Initial states:
  l[0]    s[0]    s[-1]    s[-2]    s[-3]    s[-4]    s[-5]    s[-6]    s[-7]
1807.921  27.01543 -11.25932  0.01014284 -121.9422  199.759  149.5166  37.72519  75.45785
  s[-8]    s[-9]    s[-10]   s[-11]
59.73486  47.18978 -252.1883 -211.019

sigma^2:  0.0012


      AIC      AICc      BIC
1615.746 1620.232 1657.929
```





# Hands-on

Which smoothing method would you choose if you had to choose the same method for all series? Why?



## Hands-on (#2.2)

- Apply Holt-Winter's exponential smoothing (with multiplicative seasonality) to sales of fortified wine.  
Use the training data to fit your model.

# Measuring forecasting accuracy

Forecast accuracy (or predictive accuracy) describes how closely forecasts are to the actual values. Hence, we compare the actual and forecast values and look directly at forecast errors.

$$\text{Forecast error: } e_t = y_t - F_t$$

- It is common to examine the predictive accuracy on the validation period data
- Several measures of predictive accuracy are commonly used:

Average error

$$\frac{1}{v} \sum_{i=1}^v e_t$$

MAE or MAD

(mean absolute error/deviation)

$$\frac{1}{v} \sum_{i=1}^v \left| \frac{e_t}{y_t} \right| \times 100$$

MAPE

(mean absolute percentage error)

$$\frac{1}{v} \sum_{i=1}^v |e_t|$$

RMSE

(root mean squared error)

$$\sqrt{\frac{1}{v} \sum_{i=1}^v e_t^2}$$

# Measuring forecasting accuracy

`accuracy(model estimate object)` – Gives many accuracy measurements for **training** dataset.

`accuracy(forecast object, data set object)` – Gives many accuracy measurements for the **forecasting** dataset.

```
fit.ets.HW <- train.ridership |>  
  model(ets = ETS(Ridership ~ error("M") + trend("A") + season("A")))  
fc.ets.HW <- fit.ets.HW |> forecast(h = nrow(valid.ridership))
```

```
accuracy(fit.ets.HW)
```

```
accuracy(fc.ets.HW, ridership)
```

```
> accuracy(fit.ets.HW)  
# A tibble: 1 × 10  
  .model .type    ME  RMSE  MAE    MPE  MAPE  MASE  RMSSE  ACF1  
  <chr>  <chr>  <dbl> <dbl> <dbl>  <dbl> <dbl> <dbl> <dbl> <dbl>  
1 ets    Training 0.138  56.5  45.2 -0.0792 2.58 0.548 0.569 0.0606  
> accuracy(fc.ets.HW, ridership)  
# A tibble: 1 × 10  
  .model .type    ME  RMSE  MAE    MPE  MAPE  MASE  RMSSE  ACF1  
  <chr>  <chr>  <dbl> <dbl> <dbl>  <dbl> <dbl> <dbl> <dbl> <dbl>  
1 ets    Test    33.7  76.7  62.2  1.58  3.12 0.754 0.772 0.617  
> |
```